

Your Paper Title Here

Anonymous Author(s)

ABSTRACT

Your abstract text goes here. Keep it concise, focused on the core database or data management challenge, your approach, and key experimental findings.

CCS CONCEPTS

• Information systems → Data management systems.

KEYWORDS

databases, query processing, distributed systems

1 INTRODUCTION

2 RELATED WORK

3 METHODOLOGY

3.1 Preliminaries

3.2 Query Delegation in Discovery Settings

Dynamically integrating conjunctive sub-queries during federated query processing in discovery-based querying at runtime is difficult due to the unbounded nature of the query execution. To show this, we first introduce a sufficient condition for correctly delegating conjunctive queries in the traditional federated case. Then we will show how deciding whether we can correctly delegate these queries in the discovery-based querying case is impossible when not all sources have been discovered.

First we define our discovery-based federated query setting in Definition 1 and 2.

Definition 1 (Discovery setting). A *discovery setting* is a tuple $W = \langle \mathbb{D}, \text{data}, \mathcal{D}_0, \text{disc} \rangle$: a universe of sources $\mathbb{D}$, a map $\text{data} : \mathbb{D} \rightarrow 2^T$ giving the graph each one holds, a set of seeds $\mathcal{D}_0 \subseteq \mathbb{D}$, and a *discovery function* disc assigning to every BGP B a map $\text{disc}_B : \mathbb{D} \times 2^T \rightarrow 2^{\mathbb{D}}$, where $\text{disc}_B(d, \text{data}(d))$ are the sources that a retrieved d exposes while B is evaluated.

Definition 2 (Reachable set, reference answer). Iterating discovery from the seeds, let $\mathcal{D}_{i+1} := \mathcal{D}_i \cup \bigcup_{d \in \mathcal{D}_i} \text{disc}_B(d, \text{data}(d))$. The *reachable set* is $\text{reach}_B(W) := \bigcup_{i \geq 0} \mathcal{D}_i$, and the *reference answer* to B is $\llbracket B \rrbracket_{G_{\text{reach}_B(W)}}$. Correctness of a delegation is always measured against this reference answer, never against the universe $\mathbb{D}$: sources that are never reached cannot contribute to it.

Given this definition, we recover the traditional federated query processing setting, in which all sources are known beforehand, by taking $\mathcal{D}_0 = \mathbb{D}$ and $\text{disc}_B \equiv \emptyset$.

Next we define an operator to determine the *relevance* of a source for a query execution setting. Relevance records, for each triple pattern and each source, whether that source produces matches. An example of such an operator is the usage of ASK queries to determine exclusive groups in FedX [3].

Definition 3 (Source relevance). Fix a setting W and a BGP B , and write $\mathcal{D} := \text{reach}_B(W)$ for its reachable set; all that follows is

relative to $\mathcal{D}$. Let $\mathcal{P}$ be the set of triple patterns. We write $\llbracket t \rrbracket_d$ for the evaluation of a triple pattern t over the data of source d . The *relevance indicator* $\Delta : \mathcal{P} \times \mathcal{D} \rightarrow \{0, 1\}$ is

$$\Delta(t, d) = \begin{cases} 1, & \text{if } \llbracket t \rrbracket_d \neq \emptyset, \\ 0, & \text{otherwise,} \end{cases}$$

and $\mathcal{D}_t := \{d \in \mathcal{D} \mid \Delta(t, d) = 1\}$ denotes the *relevant sources* of t .

Thus, relevance means that a data source returns answers for a given triple pattern. Next, we introduce a definition of *correct* delegation of conjunctive queries to a set of sources.

Definition 4 (Delegation). Let B be a BGP, $B' \subseteq B$ the delegated fragment, and $S \subseteq \mathcal{D}$ the (only) sources to which B' is sent as a single conjunctive request, where as in Definition 3, $\mathcal{D} = \text{reach}_B(W)$. Write $G_A := \bigcup_{d \in A} \text{data}(d)$ for $A \subseteq \mathcal{D}$. The delegated plan obtains for B'

$$R := \left(\bigcup_{d \in S} \llbracket B' \rrbracket_{\text{data}(d)} \right) \cup \llbracket B' \rrbracket_{G_{\mathcal{D} \setminus S}},$$

and the delegation is *correct* if

$$R \bowtie \llbracket B \setminus B' \rrbracket_{G_{\mathcal{D}}} = \llbracket B \rrbracket_{G_{\mathcal{D}}}.$$

This definition follows some of the definitions of federation of FedQPL [1]. Next we give a sufficient condition for correctly delegating conjunctive (sub)-queries to federated sources.

Theorem 1 (Correct delegation). Let W be a federated setting, that is $\text{disc}_B \equiv \emptyset$, so that $\mathcal{D} = \text{reach}_B(W) = \mathcal{D}_0$ is known in advance. Let $t_1, \dots, t_n$ be a set of triple patterns forming a BGP B , $B' \subseteq B$ a sub-BGP, and let $d \in \mathcal{D}$. For d we define $\mathcal{D}' = \mathcal{D} \setminus \{d\}$ the set containing all other sources. Then we can correctly delegate B' to d if $\bigcup_{t_i \in B'} \{d' \in \mathcal{D}' \mid \Delta(t_i, d') = 1\} = \emptyset$

PROOF. We prove the special case of Definition 4 with $S = \{d\}$, so that $\mathcal{D} \setminus S = \mathcal{D}'$. We assume $B' \neq \emptyset$ and that distinct sources use disjoint sets of blank nodes.

The argument rests on the fact that a *single* triple pattern distributes over a union of sources: for every t and every $A \subseteq \mathcal{D}$,

$$\llbracket t \rrbracket_{G_A} = \bigcup_{e \in A} \llbracket t \rrbracket_{\text{data}(e)}, \quad (1)$$

since matching one pattern involves no join. Note that (1) fails for BGPs, which is exactly what makes delegation non-trivial.

By hypothesis no source other than d is relevant to any pattern of B' , so Definition 3 gives

$$\llbracket t' \rrbracket_{\text{data}(d')} = \emptyset \quad \text{for all } t' \in B', d' \in \mathcal{D}'. \quad (2)$$

Step 1. A BGP is the join of its patterns, and each pattern distributes over the sources by (1):

$$\llbracket B \rrbracket_{G_{\mathcal{D}}} = \bowtie_{t \in B} \llbracket t \rrbracket_{G_{\mathcal{D}}} = \bowtie_{t \in B} \bigcup_{e \in \mathcal{D}} \llbracket t \rrbracket_{\text{data}(e)}.$$

Step 2. Since $\bowtie$ is associative and commutative, the delegated fragment splits off:

$$\llbracket B \rrbracket_{G_{\mathcal{D}}} = \left(\bowtie_{t' \in B'} \bigcup_{e \in \mathcal{D}} \llbracket t' \rrbracket_{\text{data}(e)} \right) \bowtie \llbracket B \setminus B' \rrbracket_{G_{\mathcal{D}}}.$$

Step 3. As $\mathcal{D} = \mathcal{D}' \uplus \{d\}$, the inner union splits, and by (2) its $\mathcal{D}'$ part is empty:

$$\bigcup_{e \in \mathcal{D}} \llbracket t' \rrbracket_{\text{data}(e)} = \underbrace{\bigcup_{d' \in \mathcal{D}'} \llbracket t' \rrbracket_{\text{data}(d')} \cup \llbracket t' \rrbracket_{\text{data}(d)}}_{= \emptyset} = \llbracket t' \rrbracket_{\text{data}(d)}.$$

Substituting into Step 2 and reassembling the join,

$$\llbracket B \rrbracket_{G_{\mathcal{D}}} = \llbracket B' \rrbracket_{\text{data}(d)} \bowtie \llbracket B \setminus B' \rrbracket_{G_{\mathcal{D}}}. \quad (3)$$

Step 4. It remains to identify R from Definition 4. Its second term vanishes, again by (1) and (2) and using $B' \neq \emptyset$:

$$\llbracket B' \rrbracket_{G_{\mathcal{D} \setminus S}} = \bigwedge_{t' \in B'} \bigcup_{d' \in \mathcal{D}'} \llbracket t' \rrbracket_{\text{data}(d')} = \bigwedge_{t' \in B'} \emptyset = \emptyset.$$

With $S = \{d\}$ the first term of R is the single request to d , so $R = \llbracket B' \rrbracket_{\text{data}(d)} \cup \emptyset = \llbracket B' \rrbracket_{\text{data}(d)}$.

Conclusion. Substituting R into (3),

$$R \bowtie \llbracket B \setminus B' \rrbracket_{G_{\mathcal{D}}} = \llbracket B \rrbracket_{G_{\mathcal{D}}},$$

which is the correctness condition of Definition 4. $\square$

Theorem 1 gives us a sufficient condition for delegating conjunctive sub-queries in the case where the reachable set $\mathcal{D}$ is known in advance. However, in discovery-based querying the sources are discovered at runtime, governed by some discovery process which returns during the query execution. As such, determining if we can correctly delegate depends on the *state* of the discovery process, which is defined in Definition 5.

Definition 5 (Discovery state). The state of a discovery setting $W = \langle \mathbb{D}, \text{data}, \mathcal{D}_0, \text{disc} \rangle$ at time T is the tuple

$$W(T) = \langle \mathbb{D}, \text{data}, \mathcal{D}_0, \text{disc}, \mathcal{D}(T) \rangle,$$

the setting W augmented with $\mathcal{D}(T)$, the set of sources that data has been called on by time T , subject to

$$\mathcal{D}_0 \subseteq \mathcal{D}(T) \subseteq \mathcal{D}(T') \subseteq \text{reach}_B(W) \quad \text{for all } T \leq T',$$

i.e. the seeds are discovered from the start, discovery only ever adds sources, and it never adds one outside what disc can reach. For any given $W(T)$ the discovery *frontier* is

$$F(T) = \left(\bigcup_{d_f \in \mathcal{D}(T)} \text{disc}(d_f, \text{data}(d_f)) \right) \setminus \mathcal{D}(T),$$

the sources that are discovered but do not yet have data called on them. Discovery is complete when $F(T) = \emptyset$.

Next, we introduce the concept of an observation of an engine of the discovery state in Definition 6, which denotes what the engine actually can *know* about its discovery setting at given time T .

Definition 6 (Observationally equivalent states). A discovery state $W(T)$ records the whole world, but a query engine at time T has access only to the *observation*

$$O(T) = \langle \mathcal{D}(T), \text{data}|_{\mathcal{D}(T)} \rangle,$$

the sources it has retrieved together with their contents. Two states $W_1(T)$ and $W_2(T)$ are *observationally equivalent*, written $W_1(T) \sim W_2(T)$, when $O_1(T) = O_2(T)$. Note that $F(T)$ is computable from $O(T)$ and disc: the extent of the frontier is observable, its contents are not. A property of states is *determined at T* when all states

observationally equivalent to $W(T)$ agree on it; an engine at time T can establish a property only if it is determined at T .

Theorem 2 (Correct delegation is undetermined under incomplete discovery). Let $W(T)$ be a discovery state for which $F(T) \neq \emptyset$. Then there exist a BGP B , a sub-BGP $B' \subseteq B$, a set of sources $S \subseteq \mathcal{D}(T)$, and two states $W_1(T)$ and $W_2(T)$, both observationally equivalent to $W(T)$, such that B' is correctly delegated to S in $W_1(T)$ but not in $W_2(T)$. Since the engine's input is $O(T)$ and all three states share it, no procedure can decide whether the delegation is correct.

PROOF. We construct two states $W_1(T)$ and $W_2(T)$, both observationally equivalent to $W(T)$, in which B' is delegated to S correctly and incorrectly respectively. The two agree on $\mathcal{D}_0$, disc and $\mathcal{D}(T)$, and on data over $\mathcal{D}(T)$; they differ only beyond the frontier, which is non-empty by hypothesis.

Take $\mathcal{D}(T) = S = \{d_1\}$ and

$$B' = B = \{(?s, p_1, ?o_1), (?s, p_2, ?o_2)\},$$

$$\text{data}(d_1) = \{(s_1, p_1, o_1), (s_1, p_2, o_2), (s_2, p_1, o_3)\}.$$

Fix some $d^* \in F(T)$, and let every reachable source other than d_1 and d^* hold no data. The two states are then obtained by setting

$$\text{data}(d^*) = \begin{cases} \{(s_3, p_2, o_4)\} & \text{in } W_1(T), \\ \{(s_2, p_2, o_4)\} & \text{in } W_2(T). \end{cases}$$

Since $d^* \notin \mathcal{D}(T)$, both states induce the same observation, and so $W_1(T) \sim W_2(T) \sim W(T)$ holds. The reachable set differs between the two, so we write G_i for $G_{\text{reach}_B(W_i)}$, the union of all data reachable in $W_i(T)$. As only d_1 and d^* hold data,

$$G_1 = \text{data}(d_1) \cup \{(s_3, p_2, o_4)\}, \quad G_2 = \text{data}(d_1) \cup \{(s_2, p_2, o_4)\}.$$

Write

$$\mu_1 = \{?s \mapsto s_1, ?o_1 \mapsto o_1, ?o_2 \mapsto o_2\},$$

$$\mu_2 = \{?s \mapsto s_2, ?o_1 \mapsto o_3, ?o_2 \mapsto o_4\}.$$

Neither added triple matches $(?s, p_1, ?o_1)$, so $\llbracket B' \rrbracket_{\text{data}(d^*)} = \emptyset$ in both states. By Definition 4 we therefore have

$$R_1 = R_2 = \llbracket B' \rrbracket_{\text{data}(d_1)} = \{\mu_1\}.$$

Since $B \setminus B' = \emptyset$ we have $\llbracket B \setminus B' \rrbracket_{G_i} = \{\mu_\emptyset\}$, the join identity, so the correctness condition reduces to $R_i = \llbracket B \rrbracket_{G_i}$.

In $W_1(T)$ the added triple has subject s_3 , which no p_1 -triple mentions, so it contributes no solution and

$$\llbracket B \rrbracket_{G_1} = \{\mu_1\} = R_1 :$$

the delegation is correct. In $W_2(T)$ the added triple joins with $(s_2, p_1, o_3) \in \text{data}(d_1)$ on $?s$, so

$$\llbracket B \rrbracket_{G_2} = \{\mu_1, \mu_2\} \neq R_2 :$$

the delegation is not correct.

The triple at d^* yields no solution on its own, and (s_2, p_1, o_3) yields none within d_1 ; only their join does. A single conjunctive request to d_1 cannot produce it, and neither can one to d^* , so the solution is lost precisely because the join crosses the frontier. As $W_1(T)$ and $W_2(T)$ present the engine with the same observation $O(T)$ while disagreeing on correctness, no procedure reading only $O(T)$ can decide it. $\square$

Lemma 1 (An empty frontier certifies complete discovery). $F(T) = \emptyset$ if and only if $\mathcal{D}(T) = \text{reach}_B(W)$.

PROOF. ($\Rightarrow$) The inclusion $\mathcal{D}(T) \subseteq \text{reach}_B(W)$ holds by Definition 5. For the converse, $F(T) = \emptyset$ says exactly that

$$\text{disc}(d_f, \text{data}(d_f)) \subseteq \mathcal{D}(T) \quad \text{for every } d_f \in \mathcal{D}(T),$$

i.e. that $\mathcal{D}(T)$ is closed under disc ; and $\mathcal{D}_0 \subseteq \mathcal{D}(T)$, again by Definition 5. By induction, if $\mathcal{D}_i \subseteq \mathcal{D}(T)$ then

$$\mathcal{D}_{i+1} = \mathcal{D}_i \cup \bigcup_{d \in \mathcal{D}_i} \text{disc}(d, \text{data}(d)) \subseteq \mathcal{D}(T)$$

by closure, so $\text{reach}_B(W) = \bigcup_{i \geq 0} \mathcal{D}_i \subseteq \mathcal{D}(T)$.

($\Leftarrow$) If $d \in \text{reach}_B(W)$ then $d \in \mathcal{D}_i$ for some i , hence

$$\text{disc}(d, \text{data}(d)) \subseteq \mathcal{D}_{i+1} \subseteq \text{reach}_B(W).$$

Thus $\text{reach}_B(W)$ is closed under disc , and $F(T) = \emptyset$ follows. $\square$

{TODO: Don't like the text here} is what makes the frontier useful: $F(T) = \emptyset$ is a condition the engine can test, whereas $\mathcal{D}(T) = \text{reach}_B(W)$ is a statement about a set it cannot inspect. Together with Theorem 2 it fixes exactly when delegation can be established at all.

Corollary 1 (Delegation requires complete discovery). The correctness of every delegation of a sub-BGP $B' \subseteq B$ to a set $S \subseteq \mathcal{D}(T)$ is determined at T if and only if $F(T) = \emptyset$.

PROOF. ($\Leftarrow$) Let $F(T) = \emptyset$ and let $W_1(T), W_2(T)$ be observationally equivalent to $W(T)$. As $F(T)$ is computable from $O(T)$, both have an empty frontier, so $\mathcal{D} = \text{reach}_B(W) = \mathcal{D}(T)$ in both by Lemma 1. Then $G_{\mathcal{D}} = \bigcup_{e \in \mathcal{D}(T)} \text{data}(e)$ agrees in both, and since $S \subseteq \mathcal{D}(T) = \mathcal{D}$, so do $G_{\mathcal{D} \setminus S}$ and $\text{data}(d)$ at each $d \in S$. Every term of the condition of Definition 4 therefore agrees, and with it the verdict.

($\Rightarrow$) By contraposition. If $F(T) \neq \emptyset$, then Theorem 2 supplies a B' , an S , and two states observationally equivalent to $W(T)$ that disagree on the correctness of delegating B' to S . That delegation is not determined at T . $\square$

3.3 Authority Claims

From Corollary 1, runtime federation in discovery-based querying is impossible until every source is discovered. With source counts in the thousands **{TODO: cite here}**, this is infeasible.

However, while classical query semantics require all data to be considered in query execution, this is not always an intuitive requirement. Consider a federation of product catalogues, in which a manufacturer publishes the specifications of the products it makes while resellers and comparison services republish copies of them. A query over a product's specification is answered from the manufacturer alone. However, if any copies of, different, or wrong specifications exist in the federation, they are considered equally 'true'.

We will work towards a query execution model that resolves this observation while also allowing the delegation of queries *during* query execution. First we introduce the concept of a *composite source*, an aggregating view over sources in $\mathbb{D}$ which contains no data of its own.

Definition 7 (Composite source). Fix a setting $W = \langle \mathbb{D}, \text{data}, \mathcal{D}_0, \text{disc} \rangle$. A source $c \in \mathbb{D}$ is *composite* when it holds no data of its own and is instead given by a set $\mathcal{D}_c \subseteq \mathbb{D}$ of sources it aggregates. It contains all of their data,

$$\text{data}(c) = \bigcup_{d \in \mathcal{D}_c} \text{data}(d),$$

and a request to c is evaluated over all of it. Sources that are not composite are *base* sources. We require the aggregation relation to be well-founded such that no source aggregates itself either directly or through a chain. Thus, when unfolding $\text{data}(c)$ it is expressed in base sources.

In the case $\mathcal{D}_c = \{d\}$ with d a base source, $\text{data}(c) = \text{data}(d)$. Then, c is equivalent to d .

{TODO: Add this to after claim is defined} and interchangeable when c claims (Def. 8) nothing.

Next we introduce *claims*, which declare a subset of data the composite source aggregates over that matches a certain criterion.

Definition 8 (Claim). A composite source c may *declare* a claim. A claim is a *triple-wise* criterion $\sigma_c : \mathcal{T} \rightarrow \{0, 1\}$, an indicator function mapping triples to claim membership.

A claim is published with c and thus observed once c is retrieved. The core idea on which our query execution model builds is the *principle of exclusive claims*: queries are answered as if every claim were exclusive, that is, as if no source other than c held a triple claimed by c . Like the closed-world assumption, this does not state how the federation environment is; it fixes how we answer queries over it. The data need not satisfy it, and on the Web it often does not. Definition 10 makes the principle precise by fixing the answer on any data, and the engine enforces that answer, which we discuss in Section **{TODO: ?}**

Claims indicate that a composite source is the only source for answers for triples that matched the claim indicator function. Engines can then reason about this claim to determine for what data delegating BGPs to composite sources declaring a claim is correct. As an example, in the product catalogue federation, the manufacturer declares that statements about its products are to be taken from it alone, so that copies held by resellers are disregarded.

Nothing so far prevents a source from claiming data it should not be allowed to. On the Web, any source could claim authority over, for example, DBpedia without being affiliated with it, and thereby drop every DBpedia result held elsewhere. We therefore bound each claim by what its source is entitled to.

Definition 9 (Admissibility, effective claim). An *admissibility function* assigns every composite c a criterion $\rho_c : \mathcal{T} \rightarrow \{0, 1\}$, the triples c is entitled to claim. It is part of the setting, not of any declaration, and depends only on the triple and the identity of c . The *effective claim* of c is

$$\alpha_c(t) = \sigma_c(t) \cdot \rho_c(t),$$

with $\alpha_d \equiv 0$ for base sources and for composites that declare nothing. A source d *claims* t when $\alpha_d(t) = 1$, and the principle of exclusive claims is read with effective claims.

An illegitimate claim thus shrinks to its admissible part rather than being rejected whole: a composite claiming both DBpedia and its own namespace keeps the latter.

Definition 10 (Claim semantics). A triple $t \in \text{data}(d)$ is *admitted from d* if

$$\alpha_d(t) = 1 \quad \text{or} \quad \alpha_c(t) = 0 \text{ for all } c \in \mathcal{D} \setminus \{d\}, \quad (4)$$

and for $A \subseteq \mathcal{D}$ the *claim-respecting graph* is

$$G_A^\alpha = \bigcup_{d \in A} \{t \in \text{data}(d) \mid t \text{ is admitted from } d\}.$$

The *claim semantics* of B is $\llbracket B \rrbracket_{G_D^\alpha}$. It replaces the reference answer of Definition 2, and correctness of a delegation (Definition 4) is measured with G_A^α in place of G_A .

Thus, any claimed triple only occurs within the source that claims it, while unclaimed triples are left in the data as-is. If no source claims anything, $G_A^\alpha = G_A$ and the claim semantics is the reference answer of Definition 2.

{TODO: Possibly tie to literature here about subwebs and relational?}

We keep admissibility abstract until Section 3.6, where we instantiate it by namespaces; the results in between hold for any admissibility function.

Assumption 1 (Disjoint claims). *No triple is claimed by two sources: for $c \neq c'$ and every $t \in \mathcal{T}$, $\alpha_c(t) \cdot \alpha_{c'}(t) = 0$.*

Section 3.6 gives a sufficient condition for namespace claims. We discuss relaxing the assumption in Section 3.8.1. **{TODO: What happens if we relax this?}**

3.4 Delegating Conjunctive Queries

Under this query execution model, we now show that correct delegation to composite sources can be established under *incomplete* discovery.

Lemma 2 (Claimed triples occur only in their claimant). *Under Assumption 1, let $c \in \mathcal{D}$ claim τ . Then τ is admitted from no source $d \neq c$. Consequently, for every $A \subseteq \mathcal{D}$, $\tau \in G_A^\alpha$ implies $c \in A$ and $\tau \in \text{data}(c)$.*

PROOF. Suppose τ is admitted from some $d \neq c$. The second clause of (4) fails, as $c \in \mathcal{D} \setminus \{d\}$ claims τ . So the first holds and d claims τ , contradicting Assumption 1. $\square$

A delegated composite answers over all of $\text{data}(c)$, so its answer must be filtered to what the claim semantics admits. The engine can always check the claim of c , as α_c is known once c is retrieved. For a BGP B' we therefore keep the solutions whose matched triples c all claims,

$$\Phi_c(\Omega) = \{ \mu \in \Omega \mid c \text{ claims } \mu(t) \text{ for every } t \in B' \},$$

and evaluate the rest of B' over the other sources. This refines Definition 4 for $S = \{c\}$:

$$R := \Phi_c(\llbracket B' \rrbracket_{\text{data}(c)}) \cup \llbracket B' \rrbracket_{G_{\mathcal{D} \setminus \{c\}}^\alpha}.$$

Unclaimed triples held by c are also held by the sources it aggregates, so they are found by the second term, provided these sources are reachable.

Assumption 2 (Reachable aggregation). *If a composite c is in $\mathcal{D}$, so is every source it aggregates: $\mathcal{D}_c \subseteq \mathcal{D}$.*

This holds, for instance, when a composite exposes its sources and disc follows them. What can go wrong is a solution that the two terms split between them.

Definition 11 (Claim-uniform sub-query). Let $G = G_D^\alpha$. A BGP B' is *claim-uniform* for a composite c if every $\mu \in \llbracket B' \rrbracket_G$ has c claim either all of $\mu(t)$, $t \in B'$, or none of them.

Theorem 3 (Delegation to a claim). *Under Assumptions 1 and 2, let $B' \subseteq B$ be non-empty, $c \in \mathcal{D}(T)$ a composite, and $G = G_D^\alpha$. Then $R = \llbracket B' \rrbracket_G$ if and only if B' is claim-uniform for c ; in that case delegating B' to c is correct. Consequently, whenever claim uniformity follows from B' and α_c alone, correctness is determined at T , whether or not $F(T) = \emptyset$.*

PROOF. Write $G_c := G_{\{c\}}^\alpha$ and $G_{-c} := G_{\mathcal{D} \setminus \{c\}}^\alpha$.

($R \subseteq \llbracket B' \rrbracket_G$.) A triple of $\text{data}(c)$ that c claims is admitted from c by the first clause of (4), so a solution kept by Φ_c maps into $G_c \subseteq G$. The second term maps into $G_{-c} \subseteq G$.

($\llbracket B' \rrbracket_G \subseteq R$ under uniformity.) Let $\mu \in \llbracket B' \rrbracket_G$. If c claims every $\mu(t)$, each is admitted only from c by Lemma 2, hence lies in $\text{data}(c)$; so $\mu \in \llbracket B' \rrbracket_{\text{data}(c)}$ and Φ_c keeps it. If c claims no $\mu(t)$, let $\tau = \mu(t)$ be admitted from d . If $d \neq c$, then $\tau \in G_{-c}$. If $d = c$, then τ is unclaimed, as c does not claim it and by (4) no other source does; it is held by some $d' \in \mathcal{D}_c$, which lies in $\mathcal{D}$ by Assumption 2 and admits τ , so again $\tau \in G_{-c}$. Hence $\mu \in \llbracket B' \rrbracket_{G_{-c}}$.

(Necessity.) Let $\mu \in \llbracket B' \rrbracket_G$ mix claimed and unclaimed triples, with c claiming $\mu(t_1)$ but not $\mu(t_2)$. Then Φ_c discards μ , and $\mu(t_1) \notin G_{-c}$ by Lemma 2, so $\mu \notin R$.

Given $R = \llbracket B' \rrbracket_G$, $R \bowtie \llbracket B \setminus B' \rrbracket_G = \llbracket B \rrbracket_G$, which is correctness. Finally, Φ_c depends only on α_c , and the argument quantifies over every source of $\mathcal{D}$, discovered or not; so whenever uniformity is guaranteed by B' and α_c alone, the verdict is the same in every state observationally equivalent to $W(T)$. $\square$

Theorem 3 pins down what remains of the adversary of Theorem 2: it wins only by building a solution that mixes a triple claimed by c with one that is not, so that the join crosses the boundary of c . Claim uniformity is exactly the absence of such solutions. A triple placed behind the frontier that c claims is admitted only from c , and one that c does not claim is found by the second term; neither breaks a uniform sub-query.

The simplest uniform case needs no knowledge of the data at all: if c claims every triple matching every pattern of B' , all solutions are fully claimed. We say B' is then *within the claim* of c ; the second term of R is empty and Φ_c is the identity. Under namespace claims (Section 3.6) this requires constant subjects in the namespace of c . Uniformity is much weaker: it only forbids mixed solutions, and Section 3.6.2 shows that subject-determined claims guarantee it for star-shaped sub-queries regardless of the data. A composite that claims nothing has $\Phi_c(\Omega) = \emptyset$, so delegating to it is vacuous: we delegate only to claimants, and a single document only as a degenerate composite that declares a claim.

{TODO: Lemma on how theorem 1 holds if we filter any sources in scope of c both backwards and forwards in time. Mention how the filtering changes nothing given the assumptions / semantics of our model}

PROOF. Sketch Part 1. If we filter $\rightarrow$ Theorem 1 can be directly applied Part 2. Filtering is possible: forward $\rightarrow$ easy just take $\text{disc}()$ and apply setminus of in scope documents. Backwards $\rightarrow$ deduplication, we only have duplicate results. So filter any intermediate results at time T in $\text{data}(\text{scope})$ and filter any produced results in composite resource on B prime $\square$

3.5 Monotonicity

In classical discovery-based querying, a result is sound the moment it is produced: as $\mathcal{D}(T)$ grows, so does $G_{\mathcal{D}(T)}$, and BGP evaluation is monotone. Claims break this, as discovering a composite can disregard triples that were already used. We make this precise for the answer an engine can compute at time T , which can only take the claims of retrieved composites into account.

Definition 12 (Answer at time T). For a discovery state $W(T)$, with claims under α , let

$$G^T = \bigcup_{d \in \mathcal{D}(T)} \{t \in \text{data}(d) \mid d \text{ claims } t \text{ or no } c \in \mathcal{D}(T) \text{ claims } t\},$$

and $\Omega(T) = \llbracket B \rrbracket_{G^T}$. The answer is *monotone* if $\Omega(T) \subseteq \Omega(T')$ whenever $T \leq T'$.

Once discovery is complete, $\mathcal{D}(T) = \mathcal{D}$ by Lemma 1 and $G^T = G_{\mathcal{D}}^{\alpha}$, so $\Omega(T)$ is the answer under claims. Without claims, $G^T = G_{\mathcal{D}(T)}$ and Ω is monotone.

Proposition 1 (Claims break monotonicity). *There is a setting with times $T \leq T'$ such that $\Omega(T) \not\subseteq \Omega(T')$.*

PROOF. Take $B = \{(s, p, ?o)\}$ and a composite c that claims $\tau = (s, p, o)$. Let a base source $d \notin \mathcal{D}_c$ hold τ while $\tau \notin \text{data}(c)$, and let d be retrieved before c . If $\mathcal{D}(T) = \{d\}$, then $\mu = \{s \mapsto s, ?o \mapsto o\} \in \Omega(T)$. Once $c \in \mathcal{D}(T')$, c claims τ , so τ is no longer admitted from d , nor from c as $\tau \notin \text{data}(c)$. Hence $\mu \notin \Omega(T')$. $\square$

The triple τ is claimed by c but held only elsewhere; under namespace claims it is a statement by d about a resource in the namespace of c . It is exactly the kind of triple the adversary of Theorem 2 places behind the frontier, and exactly what Theorem 3 needs disregarded. Non-monotonicity is thus the price of delegation under incomplete discovery rather than an artefact of our definitions. **{TODO: Formalise as a necessity result?}** It is, however, confined to such triples.

Definition 13 (Covered triple). A triple t is *covered* if every composite $c \in \mathcal{D}$ that claims t also holds it, $t \in \text{data}(c)$.

Lemma 3 (Covered solutions are stable). *Under Assumption 1, let $\mu \in \Omega(T)$ be such that $\mu(t)$ is covered for every $t \in B$. Then $\mu \in \Omega(T')$ for every $T' \geq T$, and $\mu \in \llbracket B \rrbracket_{G_{\mathcal{D}}^{\alpha}}$.*

PROOF. Let $t \in B$ and $\tau = \mu(t) \in G^T$, so some $d \in \mathcal{D}(T)$ holds τ . By Assumption 1 at most one source claims τ . If no source of $\mathcal{D}(T')$ claims it, τ is admitted at T' from $d \in \mathcal{D}(T')$. Otherwise its claimant $c \in \mathcal{D}(T')$ holds τ , as τ is covered, and τ is admitted from c . Hence $\mu(B) \subseteq G^{T'}$, and $\mu \in \Omega(T')$ as $\text{dom}(\mu) = \text{var}(B)$. The same argument with $\mathcal{D}$ in place of $\mathcal{D}(T')$ gives $\mu \in \llbracket B \rrbracket_{G_{\mathcal{D}}^{\alpha}}$. $\square$

Which triples are covered depends on admissibility. For namespace claims it follows from where a triple is stored, which Section 3.6 uses to resolve claim status in practice.

3.6 Subject-Namespace Claims

The results so far hold for any admissibility function satisfying Assumption 1. Two further questions need more structure: when a sub-query whose patterns have variable subjects can be delegated, and how an engine settles claim status in practice. Both rely on claims that are decided by the subject of a triple alone.

Definition 14 (Subject-determined claims). Claims are *subject-determined* if for every composite c there is a set of terms S_c such that $\alpha_c(t) = 1$ iff $\text{subj}(t) \in S_c$.

For this work we instantiate admissibility by namespaces, as discovery-based querying on the Web identifies sources by IRI. With $\text{ns}(c)$ the namespace of the IRI of c ,

$$\rho_c(t) = 1 \iff \text{subj}(t) \in \text{ns}(c),$$


and a composite that declares claims its whole namespace, $\sigma_c = \rho_c$. Claims are then subject-determined with $S_c = \text{ns}(c)$. A composite with namespace $\text{http://example.org/}$ thus claims every triple whose subject starts with it. ρ_c depends only on the identity of c , so the engine knows it as soon as it discovers c , whatever lies behind the frontier. Admissibility bounds what c may *claim*, not what it may *aggregate*; for the latter, write $\text{iri}(d)$ for the IRI of a source d .

Assumption 3 (Namespace aggregation). *A composite aggregates every source in its namespace: if $\text{iri}(d) \in \text{ns}(c)$ then $d \in \mathcal{D}_c$. {TODO: align with subweb paper}*

The picture is that of horizontal partitioning, with each composite the sole source of record for a region of the subject space. The difference is that the partitioning is declared by the participants rather than imposed by a coordinator, so no global catalogue of it exists and both the participants and the partitioning are discovered during execution. Assumption 1 holds when no composite's namespace extends another's. Disjoint aggregation alone does not suffice: a composite on $\text{http://example.org/}$ and one on $\text{http://example.org/a/}$ may aggregate disjoint documents and still both claim a triple about $\text{http://example.org/a/x}$.

3.6.1 Resolving claim status. Call $t \in \text{data}(d)$ *local* if no composite namespace separates its subject from its source: for every composite c , $\text{subj}(t) \in \text{ns}(c)$ implies $\text{iri}(d) \in \text{ns}(c)$. Under Assumption 3 a local triple is covered, since a composite claiming it has $\text{subj}(t) \in \text{ns}(c)$, hence $d \in \mathcal{D}_c$ and $t \in \text{data}(c)$. By Lemma 3, solutions built from local triples are sound as soon as they are produced; only solutions using a non-local triple, a statement about a resource in another namespace, can be invalidated by later discovery.

Remark 1 (Resolving claim status by request). In practice the engine need not wait for discovery to settle whether a non-local triple t is claimed. By Assumption 1 at most one composite can claim t , and under namespace admissibility its location is fixed by $\text{subj}(t)$ alone: it is the composite, if any, whose namespace contains $\text{subj}(t)$. The engine can therefore resolve the claim status of t with a request to this possible aggregator, for example by looking up an advertised composite along the path of $\text{subj}(t)$ and, if one claims t , asking it whether it holds t . Including these lookups in disc makes the aggregator part of $\text{reach}_B(W)$, so the outcome agrees with the answer under claims. An engine that resolves the status of every non-local triple before using it emits only solutions that remain valid, at the



(a) Static plan: t_1 and t_2 lie in different subtrees, so c cannot be added to the plan.
 (b) Eddy: routing history split t_3 into t'_3 and t''_3 across the joins, so the $t_1 t_2$ tuples of c probe only t'_3 .

Figure 1: **{TODO: caption}**

cost of one lookup per foreign namespace and one membership check per non-local triple, rather than a wait for the frontier to empty. This mirrors authoritative reasoning, where authority is decided per term by dereferencing [2]. **{TODO: Mechanism for advertising a composite, and for establishing that none exists: storage description, link header?}**

3.6.2 Star- and chain-shaped sub-queries. Subject-determined claims make the claim status of a matched triple a function of the binding of its subject. For a star whose patterns share the subject variable $?s$, all triples matched by a solution μ have subject $\mu(?s)$, so either c claims all of them or none. The star is therefore claim-uniform for every composite, whatever the data, and Theorem 3 applies at any time T ; Φ_c keeps the solutions with $\mu(?s) \in S_c$. Patterns with distinct subject variables, as in chains, admit mixed solutions. **{TODO: Formalise for stars and size-2 chains}**
{TODO: Section indicating how anchor variables are required to correctly filter out partial results} **{TODO: This then makes stars and size 2 chains attractive candidates to reduce client work, while longer chains can't really be pruned.}**

3.7 Adaptive Query Planning with Composite Sources

Given the capability of delegating these sub queries, we must now move to the system-side of our problem. Specifically, due to the dynamic discovery of composite sources and the uncertainty of capabilities of these sources, using a static query plan is difficult.

{TODO: Text referencing Figure 1}

{TODO: Show that even if we use many alternative plans we can still run into issues with overlapping composite sources}

{TODO: Stems!}

3.8 Correctness

3.8.1 Overlapping Sources. Correctness when doing overlapping sources is difficult blabla, cases: Overlapping domain, overlapping triple patterns, both, neither. Given the time dimension, how do we decide what is possible

3.9 Discovery

4 SYSTEM APPLICATION

To evaluate the performance impact of delegating sub-queries we implement hybrid federated link traversal-based query processing. Specifically, within solid blabla

Note this is a concrete application of the stuff introduced here, which is not limited to Solid and link traversal.

5 EVALUATION

TODO: Ensure correctness of multiple composite resources!! Experiments: - Full benchmark and ablations - Client-scaling test (find server-side usage and client-side scaling behavior) - Test query with increasing pod depth to find answer - Query where linear is fastest to delegate with single pod (show star-only fails and cost-based fails) - Query where linear is fastest to delegate with multiple pods (show if cost-based doesn't fail) - Query where star is fastest multiple pods (show if cost-based doesn't fail) - Multiple vs Single pod query

6 CONCLUSION

Summarize results and point to future work.

REFERENCES

- [1] Sijin Cheng and Olaf Hartig. 2020. FedQPL: a language for logical query plans over heterogeneous federations of RDF data sources. In *Proceedings of the 22nd International Conference on Information Integration and Web-based Applications & Services*. 436–445.
- [2] Aidan Hogan, Andreas Harth, and Axel Polleres. 2009. Scalable Authoritative OWL Reasoning for the Web. *International Journal on Semantic Web and Information Systems* 5, 2 (2009), 49–90. doi:10.4018/jswis.2009040103
- [3] Andreas Schwarte, Peter Haase, Katja Hose, Ralf Schenkel, and Michael Schmidt. 2011. FedX: Optimization Techniques for Federated Query Processing on Linked Data. In *The Semantic Web – ISWC 2011 (LNCS, Vol. 7031)*. Springer, 601–616. doi:10.1007/978-3-642-25073-6_38